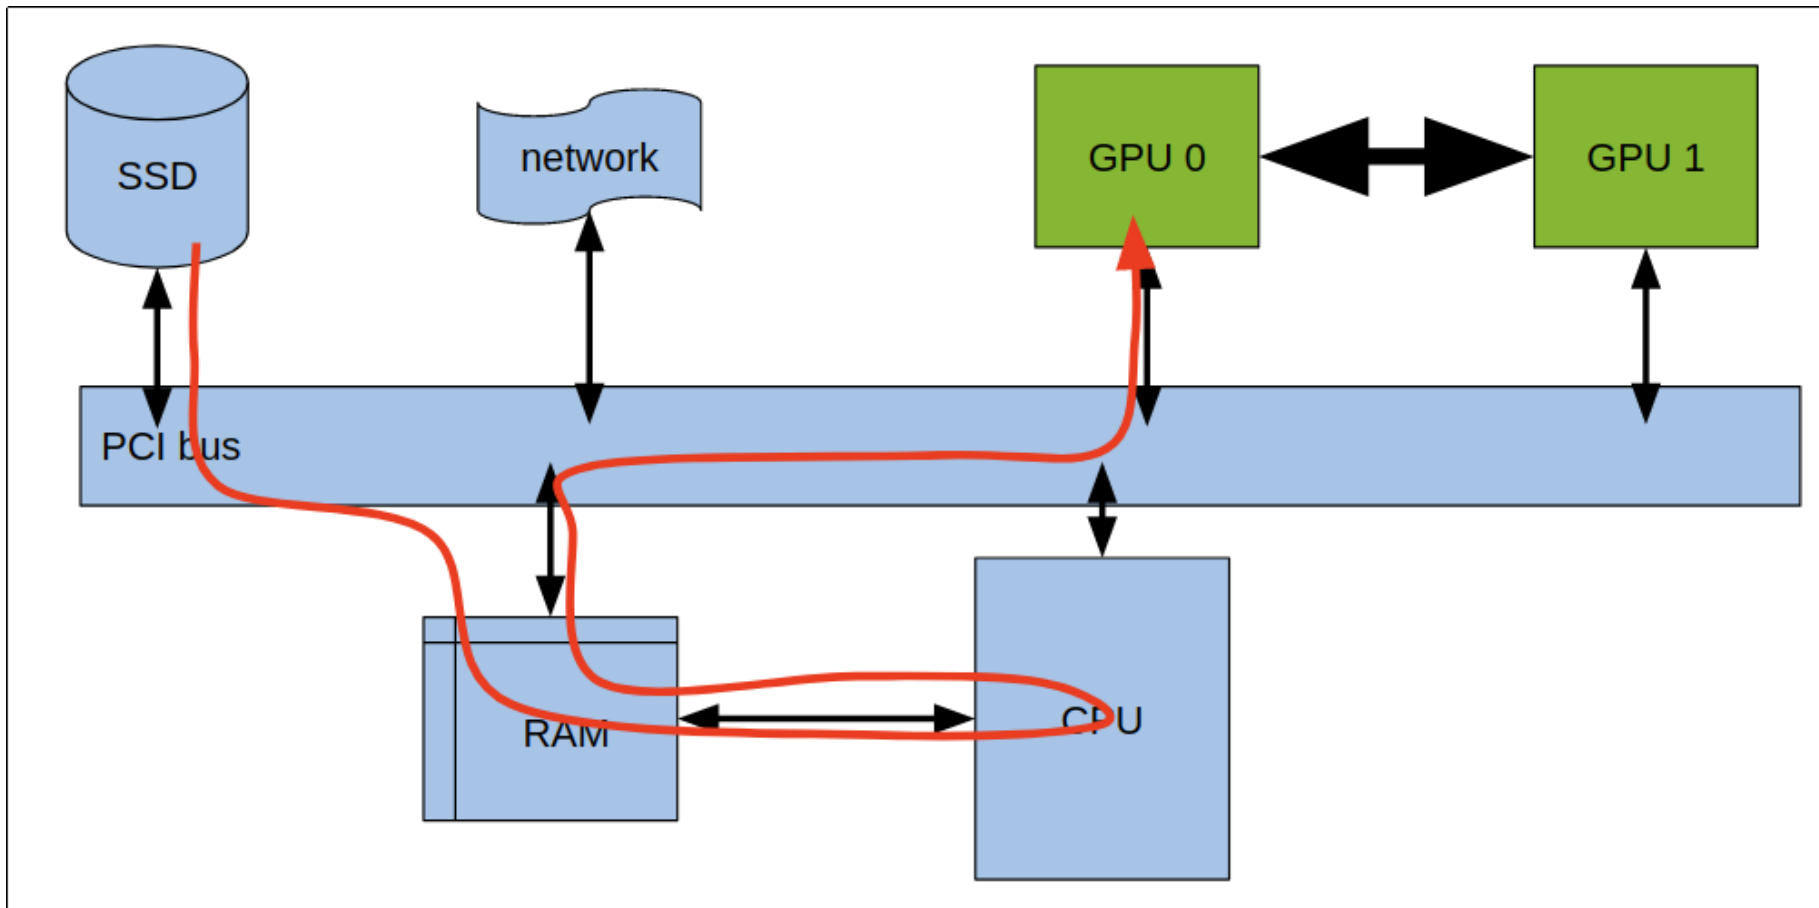


# Performance and Profiling

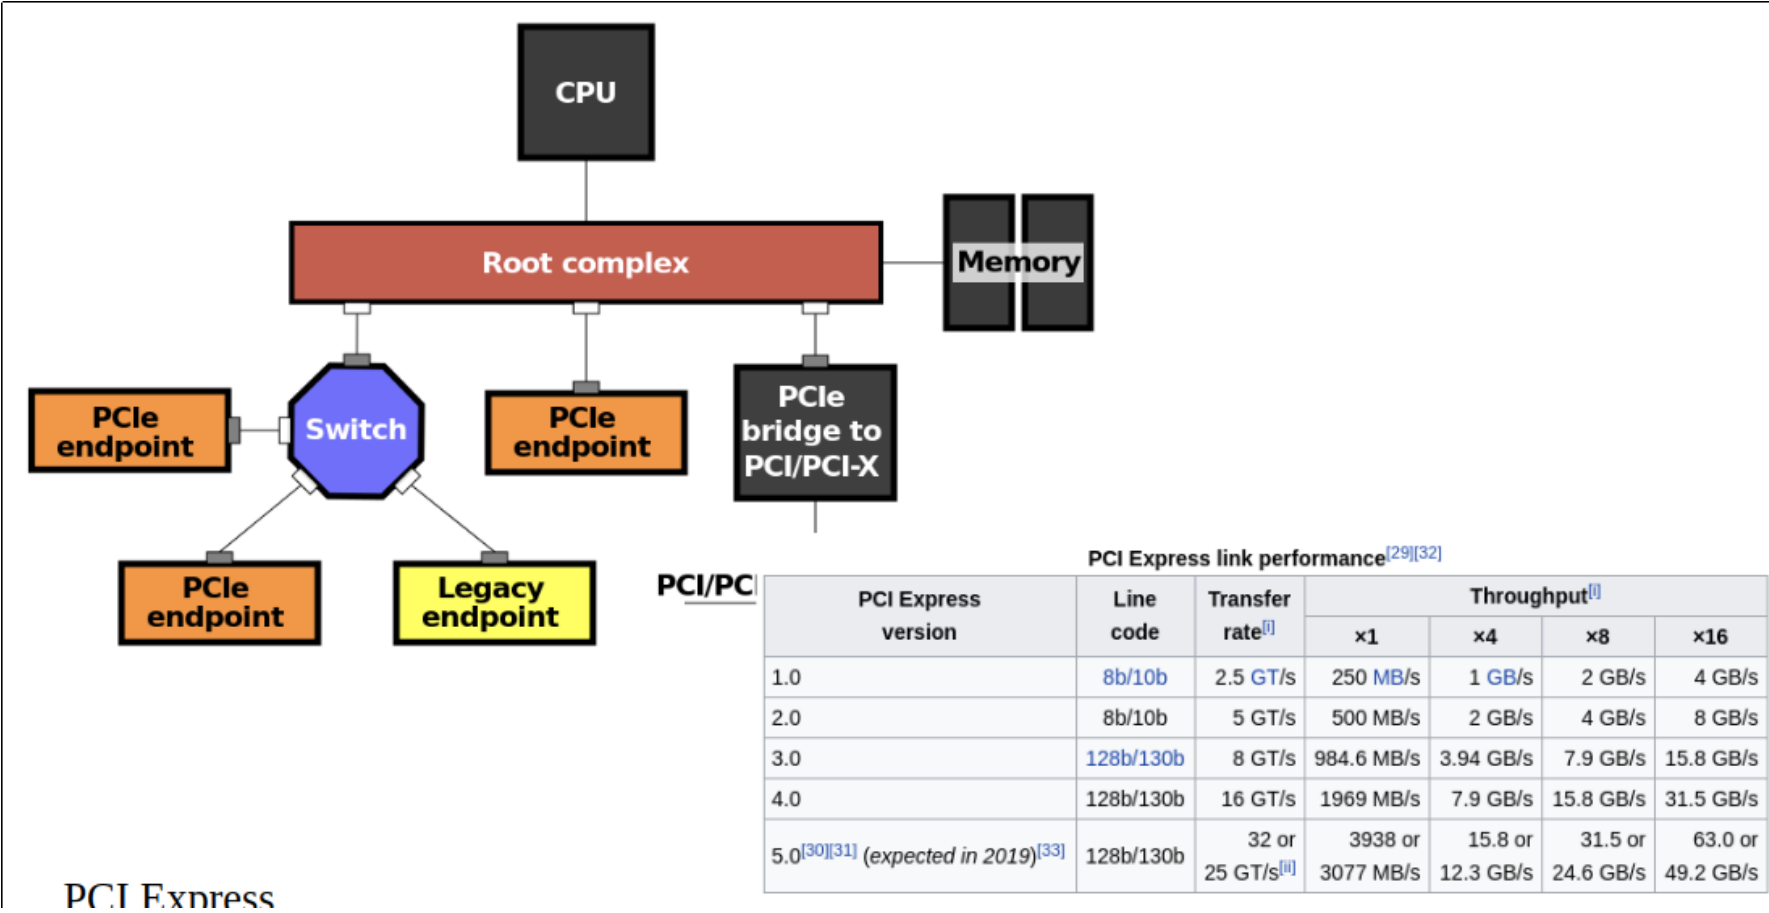
# Flow of Data during Training



# Common System Bandwidths

Hardware Component	speed	1 PB = ...
Intel i7 5930K (40 PCIe lanes)	30 GB/s	9 hours
NVLINK	40 GB/s	7 hours
PCIe x8 (GPU)	5 GB/s	2 days
SATA Interface	6 GB/s	2 days
56 Gbps Mellanox	6.8 GB/s	2 days
10 Gbps Ethernet	1 GB/s	2 weeks

# Detailed PCI Bus Structure



# Common Ways of Speeding Up Single GPU Jobs

- CPU
  - use multicore/multithreading
  - use pre-augmented data
  - avoid data copies
- GPU
  - use pinned memory for CPU/GPU transfers
  - overlay memory transfers and computation
  - switch to FP16 and use tensor cores
- speed up I/O
  - load your entire dataset into CPU/GPU memory
  - use sequential reads or NVMe

# How do you find out what to do?

- system monitoring tools (CPU, GPU, disk, network)
- manual instrumentation, logging, and performance measurements
- mock loaders / trainers
- performance analysis and visualization tools

(In rough order from easy to hard.)

# Manual Instrumentation

```
while True:
    t1 = time.time()
    inputs, targets = next(source)

    t2 = time.time()
    optimizer.zero_grad()
    outputs = model(inputs)
    loss = lossfn(outputs, targets)
    loss.backward()
    optimizer.step()
    t3 = time.time()

    loading_time = moving_average(loading_time, t2-t1)
    training_time = moving_average(training_time, t3-t2)
```

I/O should overlap with training, and  $t2-t1$  should be much smaller than  $t3-t2$

# Performance Testing

To measure limit of training performance:

- mock up the loader
- store a single batch in CPU or GPU memory

To measure limit of I/O performance:

- mock up the training = measure loading performance
- just discard each batch after loading

I/O sample rate should be higher than training sample rate.



# GPU Utilization

First thing to look at: **What is my GPU utilization?**

```
$ nvidia-smi
```

Check for:

- utilization of each GPU (should be close to 100%)
- GPU memory utilization (stay away from max)
- active processes and their usage (as few as possible)
- temperature (make sure you're cooled enough)

Also: [NVIDIA GPU profiling tools](#)



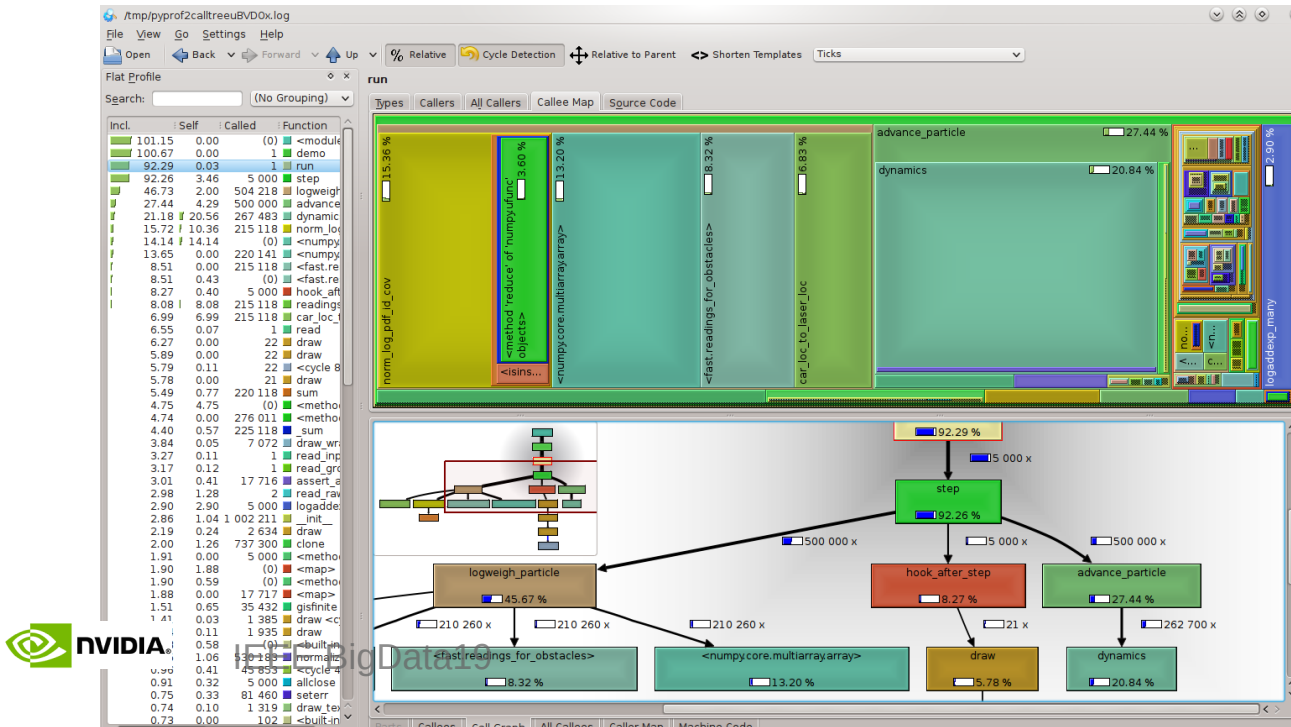
# Tools: CPU Utilization

First thing to look at: **How are my CPU cores being used?**

```
$ htop
```

Many more tools, e.g.:

```
$ python -m cProfile -o myscript.cprof myscript.py  
$ pyprof2calltree -k -i myscript.cprof
```



# Tools: PyTorch Profiler

```
import torch
import torchvision.models as models

model = models.densenet121(pretrained=True)
x = torch.randn((16, 3, 224, 224), requires_grad=True)

with torch.autograd.profiler.profile(use_cuda=True) as prof:
    model(x)

print(prof)

prof.export_chrome_trace("mytrace") # open with chrome://tracing
```

# Tools: PyTorch Profiler (Sample Output)

Name	CPU time	CUDA time	Calls	CPU total	CUDA total
conv2d	9976.544us	9972.736us	1	9976.544us	9972.736us
convolution	9958.778us	9958.400us	1	9958.778us	9958.400us
_convolution	9946.712us	9947.136us	1	9946.712us	9947.136us
contiguous	6.692us	6.976us	1	6.692us	6.976us
empty	11.927us	12.032us	1	11.927us	12.032us
mkldnn_convolution	9880.452us	9889.792us	1	9880.452us	9889.792us
batch_norm	1214.791us	1213.440us	1	1214.791us	1213.440us
native_batch_norm	1190.496us	1193.056us	1	1190.496us	1193.056us
threshold_	158.258us	159.584us	1	158.258us	159.584us
max_pool2d_with_indices	28837.682us	28836.834us	1	28837.682us	28836.834us
max_pool2d_with_indices_forward	28813.804us	28822.530us	1	28813.804us	28822.530us
batch_norm	1780.373us	1778.690us	1	1780.373us	1778.690us
native_batch_norm	1756.774us	1759.327us	1	1756.774us	1759.327us
threshold_	64.665us	66.368us	1	64.665us	66.368us
conv2d	6103.544us	6102.142us	1	6103.544us	6102.142us
convolution	6089.946us	6089.600us	1	6089.946us	6089.600us
_convolution	6076.506us	6076.416us	1	6076.506us	6076.416us
contiguous	7.306us	7.938us	1	7.306us	7.938us
empty	9.037us	8.194us	1	9.037us	8.194us
mkldnn_convolution	6015.653us	6021.408us	1	6015.653us	6021.408us
batch_norm	700.129us	699.394us	1	700.129us	699.394us

# Speedup: FP16 Computations

Important:

- intrinsically faster
- enables the use of TensorCores
- changes numerical results and can't be used with all computations

Speedups:

- between 1.5x and 5x for common models

# Speedup: Converting to FP16

## EXAMPLE

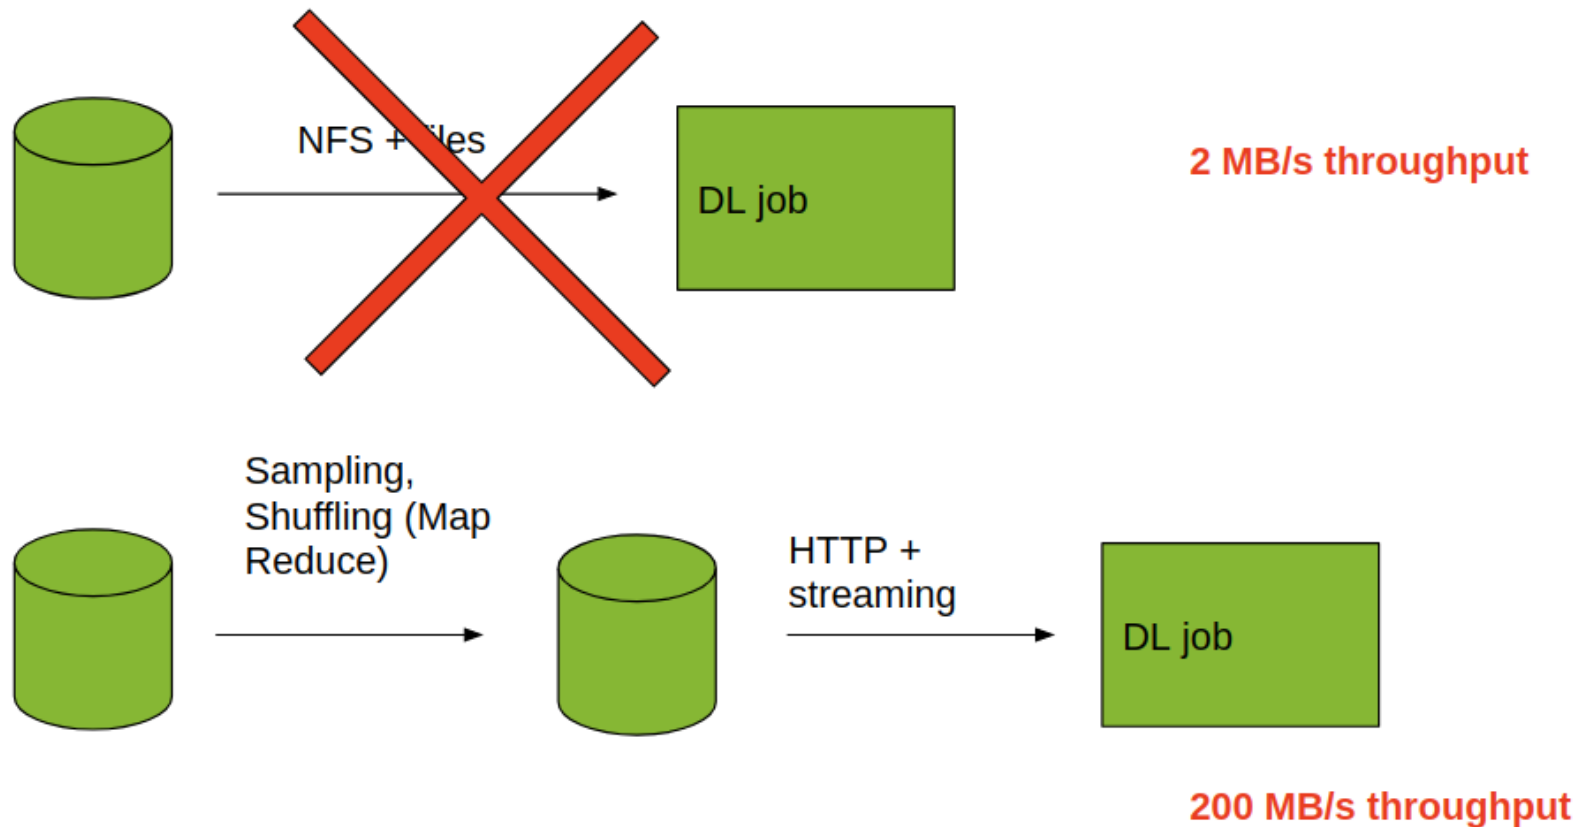
```
N, D_in, D_out = 64, 1024, 512
x = torch.randn(N, D_in, device="cuda")
y = torch.randn(N, D_out, device="cuda")

model = torch.nn.Linear(D_in, D_out).cuda()
optimizer = torch.optim.SGD(model.parameters(), lr=1e-3)
model, optimizer = amp.initialize(model, optimizer, opt_level="O1")

for t in range(500):
    y_pred = model(x)
    loss = torch.nn.functional.mse_loss(y_pred, y)

    optimizer.zero_grad()
    with amp.scale_loss(loss, optimizer) as scaled_loss:
        scaled_loss.backward()
    optimizer.step()
```

# Speedup: Sequential I/O instead of Random Access



# Linux Monitoring

Locally:

- htop, ... = watch processes
- iotop, ... = watch I/O
- nettop, ... = watch network I/O

Distributed:

- sensors in nodes/containers
- logging in log server
- visualization frontend



# Recommendations

Max Out the Expensive Stuff:

- ensure that you are getting 90%+ GPU utilization for each GPU
- check your I/O bandwidth; it should be about 150 MB/s for disks, 3000 MB/s for NVMe

Avoid Maxing Out:

- GPU memory: this will kill your job
- CPU memory: limits GPU performance
- CPU utilization: limits GPU performance
- network bandwidth: limits I/O performance (for distributed I/O and SGD)

# Some Options

- parallelize your model
- parallelize your model differently
- change random access I/O to sequential I/O
- use more CPU cores per GPU and use more I/O workers per GPU
- move I/O to separate node
- move data augmentation to separate node(s)
- use RDMA

(We will cover these later.)

# FP16 Notebook

(notebook)