

Kubernetes

Kubernetes

- Docker provides individual containers on a local machine
- Kubernetes manages collections of running containers across a cluster/datacenter
- also provides networking, storage, monitoring, service discovery

The Cluster

A cluster with 6 CPU nodes and 8 GPU nodes (running on Google GCE).

In [1]:

```
# make sure we have a running cluster
kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
gke-tmb-cluster-default-pool-7a7b0dec-71mc	Ready	<none>	11h	v1.13.11-gke.14
gke-tmb-cluster-default-pool-7a7b0dec-8xm3	Ready	<none>	11h	v1.13.11-gke.14
gke-tmb-cluster-default-pool-7a7b0dec-fdkw	Ready	<none>	11h	v1.13.11-gke.14
gke-tmb-cluster-default-pool-7a7b0dec-m2b7	Ready	<none>	11h	v1.13.11-gke.14
gke-tmb-cluster-default-pool-7a7b0dec-whjj	Ready	<none>	11h	v1.13.11-gke.14
gke-tmb-cluster-default-pool-7a7b0dec-wzq1	Ready	<none>	11h	v1.13.11-gke.14
gke-tmb-cluster-gpus-6c20b4bb-9w8f	Ready	<none>	11h	v1.13.11-gke.14
gke-tmb-cluster-gpus-6c20b4bb-bm56	Ready	<none>	11h	v1.13.11-gke.14
gke-tmb-cluster-gpus-6c20b4bb-mfrs	Ready	<none>	11h	v1.13.11-gke.14
gke-tmb-cluster-gpus-6c20b4bb-pzm4	Ready	<none>	11h	v1.13.11-gke.14
gke-tmb-cluster-gpus-6c20b4bb-r7tc	Ready	<none>	11h	v1.13.11-gke.14
gke-tmb-cluster-gpus-6c20b4bb-t4r9	Ready	<none>	11h	v1.13.11-gke.14
gke-tmb-cluster-gpus-6c20b4bb-tx26	Ready	<none>	11h	v1.13.11-gke.14
gke-tmb-cluster-gpus-6c20b4bb-zjp1	Ready	<none>	11h	v1.13.11-gke.14

In [2]:

```
kubectl delete jobs --all
kubectl delete pods --all
```

No resources found
pod "myjob-z6f49" deleted

Pods

- Kubernetes groups containers into pods
- (Docker container = whale, Pod = group of whales)
- specifications are written in YAML or JSON

In [5]:

```
kubectl delete pod/mypod || true
kubectl apply -f - <<'EOF'
apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  containers:
  - name: mypod
    image: gcr.io/research-191823/bigdata19
    command: ["nvidia-smi"]
    resources:
      limits:
        nvidia.com/gpu: "1"
  restartPolicy: Never
EOF
```

Error from server (NotFound): pods "mypod" not found
pod/mypod created

Pod Status and Logs

The Kubernetes runtime keeps track of pod status and logs.

In [6]:

```
kubect  get pods
```

NAME	READY	STATUS	RESTARTS	AGE
mypod	0/1	ContainerCreating	0	0s

In [7]:

```
sleep 15
```

In [8]:

```
kubect  logs pod/mypod
```

Mon Dec 9 16:33:29 2019

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+																	
NVIDIA-SMI 418.67		Driver Version: 418.67				CUDA Version: 10.1											
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+																	
GPU Name		Persistence-M		Bus-Id		Disp.A		Volatile Uncorr. ECC									
Fan Temp Perf		Pwr:Usage/Cap		Memory-Usage		GPU-Util		Compute M.									
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+																	
0 Tesla T4		Off		00000000:00:04.0		Off		0									
N/A 40C P8		10W / 70W		0MiB / 15079MiB		0%		Default									
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+																	
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+																	
Processes:								GPU Memory									
GPU PID Type Process name								Usage									
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+																	
No running processes found																	
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+																	

Debugging Pod Startup Problems

Sometimes pods don't get scheduled (never start running). Here are some tricks to debug this.

In [9]:

```
# sometimes pods don't schedule; there is tons of info
kubect  describe pod/mypod | sed 10q
kubect  describe pod/mypod | echo ... ${wc -l} ...
```

Name:	mypod
Namespace:	default
Priority:	0
PriorityClassName:	<none>
Node:	gke-tmb-cluster-gpus-6c20b4bb-9w8f/10.138.0.75
Start Time:	Mon, 09 Dec 2019 08:33:28 -0800
Labels:	<none>
Annotations:	kubect�.kubernetes.io/last-applied-configuration: {"apiVersion":"v1","kind":"Pod","metadata":{"annotations":{},"name":"mypod"},"namespace":"default"},"spec":{"containers":[{"command":["nvid... kubernetes.io/limit-ranger: LimitRanger plugin set: cpu request for container my pod ... 60 ...

In [10]:

```
# the Events: section usually tells you why a job didn't get assigned to a node

kubect  describe pod/mypod | grep -A100 Events:
```

Events:				
Type	Reason	Age	From	Message
----	-----	----	-----	-----
Normal	Scheduled	17s	default-scheduler	Successfully assigned default
t/mypod	to gke-tmb-cluster-gpus-6c20b4bb-9w8f			
Normal	Pulling	16s	kubelet, gke-tmb-cluster-gpus-6c20b4bb-9w8f	pulling image "gcr.io/resear
ch-191823/bigdata19"				
Normal	Pulled	16s	kubelet, gke-tmb-cluster-gpus-6c20b4bb-9w8f	Successfully pulled image "g
cr.io/research-191823/bigdata19"				
Normal	Created	16s	kubelet, gke-tmb-cluster-gpus-6c20b4bb-9w8f	Created container
Normal	Started	16s	kubelet, gke-tmb-cluster-gpus-6c20b4bb-9w8f	Started container

In [11]:

```
# nodes also have descriptions (even longer)

node=$(kubectl get nodes | awk '/gpu/{print $1; exit}')
kubectl describe node/$node | sed 10q
kubectl describe node/$node | echo ... $(wc -l) ...
```

Name: gke-tmb-cluster-gpus-6c20b4bb-9w8f
Roles: <none>
Labels: beta.kubernetes.io/arch=amd64
beta.kubernetes.io/fluentd-ds-ready=true
beta.kubernetes.io/instance-type=n1-standard-16
beta.kubernetes.io/os=linux
cloud.google.com/gke-accelerator=nvidia-tesla-t4
cloud.google.com/gke-nodepool=gpus
cloud.google.com/gke-os-distribution=cos
failure-domain.beta.kubernetes.io/region=us-west1
... 85 ...

In [12]:

```
# you want to make sure that nodes have the right allocatable resources
kubectl describe node/$node | grep -A10 Allocatable:
```

Allocatable:
attachable-volumes-gce-pd: 127
cpu: 15890m
ephemeral-storage: 47093746742
hugepages-2Mi: 0
memory: 56288600Ki
nvidia.com/gpu: 1
pods: 110
System Info:
Machine ID: 265593ea8efdb186402965bc1163ba81
System UUID: 265593EA-8EFD-B186-4029-65BC1163BA81

In [13]:

```
# also make sure there are resources available
kubectl describe node/$node | grep -A10 Allocated
```

Allocated resources:
(Total limits may be over 100 percent, i.e., overcommitted.)
Resource Requests Limits

cpu 400m (2%) 1050m (6%)
memory 210Mi (0%) 510Mi (0%)
ephemeral-storage 0 (0%) 0 (0%)
attachable-volumes-gce-pd 0 0
nvidia.com/gpu 0 0
Events: <none>

In [15]:

```
# nodes can be prevented from scheduling jobs by "taints"
kubectl describe node/$node | grep -A2 Taints:
```

Taints: nvidia.com/gpu=present:NoSchedule
Unschedulable: false
Conditions:

In [14]:

```
# only pods that tolerate the taints are scheduled
kubectl describe pod/mypod | grep -A5 Tolerations:
```

Tolerations: node.kubernetes.io/not-ready:NoExecute for 300s
node.kubernetes.io/unreachable:NoExecute for 300s
nvidia.com/gpu:NoSchedule
Events:
Type Reason Age From Message

In [16]:

```
kubectl delete pods --all

pod "mypod" deleted
```

Jobs

Jobs are like batch queuing. Job specs are a wrapper around pod specs.

In [17]:

```
kubectl delete job/myjob || true
kubectl apply -f - <<'EOF'
apiVersion: batch/v1
kind: Job
metadata:
  name: myjob
  labels:
    app: bigdata19
spec:
  backoffLimit: 0
  template:
    # below is a regular Pod spec
    spec:
      containers:
        - name: myjob
          image: gcr.io/research-191823/bigdata19
          command:
            - "/bin/bash"
            - "-c"
            - |
              nvidia-smi
          stdin: true
          tty: true
          resources:
            limits:
              nvidia.com/gpu: "1"
          restartPolicy: Never
EOF
```

Error from server (NotFound): jobs.batch "myjob" not found
job.batch/myjob created

In [18]:

```
sleep 15
```

In [19]:

```
kubectl logs job.batch/myjob
```

```

Mon Dec 9 16:33:53 2019
+-----+
| NVIDIA-SMI 418.67           Driver Version: 418.67           CUDA Version: 10.2           |
+-----+-----+
| GPU  Name                Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp   Perf          Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
+-----+-----+
|  0   Tesla T4              Off      | 00000000:00:04.0 Off |             0         |
| N/A   43C     P0             25W / 70W |  0MiB / 15079MiB |           5%      Default |
+-----+-----+

+-----+-----+
| Processes:                                     GPU Memory |
|  GPU           PID    Type   Process name                     Usage      |
+-----+-----+
| No running processes found                  |
+-----+-----+

```

In [20]:

```
kubectl get jobs
```

NAME	COMPLETIONS	DURATION	AGE
myjob	1/1	3s	16s

In [21]:

```
kubectl delete jobs --all
```

job.batch "myjob" deleted

Configmaps

In [22]:

```
# configmaps are little mountable file systems, for config information and scripts
# we put our Python scripts there
kubectl delete configmap files || true
kubectl create configmap files \
--from-file=training.py=training.py \
--from-file=helpers.py=helpers.py
```

configmap "files" deleted
configmap/files created

Running a Training Job

In [23]:

```
# with the scripts transferred, let's run actual training
# note the use of multi-line quoting for the shell script
kubectl delete job/myjob || true
kubectl apply -f - <<'EOF'
apiVersion: batch/v1
kind: Job
metadata:
  name: myjob
  labels:
    app: bigdata19
spec:
  backoffLimit: 0
  template:
    spec:
      containers:
        - name: myjob
          image: gcr.io/research-191823/bigdata19
          command:
            - "/bin/bash"
            - "-c"
            - |
              cp /files/*.py .
              python3 training.py
          stdin: true
          tty: true
          resources:
            limits:
              nvidia.com/gpu: "1"
          volumeMounts:
            - mountPath: /files
              name: files
      restartPolicy: Never
      volumes:
        - configMap:
            name: files
          name: files
EOF
```

Error from server (NotFound): jobs.batch "myjob" not found
job.batch/myjob created

Training Job

In [24]:

```
kubectl get jobs
```

NAME	COMPLETIONS	DURATION	AGE
myjob	0/1	1s	1s

In [25]:

```
sleep 30
```

In [26]:

```
kubectll logs job/myjob
```

```
/opt/conda/lib/python3.6/site-packages/torchvision/io/_video_opt.py:17: UserWarning: video reader based on ffmpeg c++ ops not available
  warnings.warn("video reader based on ffmpeg c++ ops not available")
Mon Dec 9 16:34:14 UTC 2019; myjob-fgflt; root; /workspace; GPU 0: Tesla T4 (UUID: GPU-7ffd1122-5cce-b8a1-db96-99d9e51ebbc8);
creating resnet50
    0 bs    128 per sample loss 5.57e-02 loading 8.92e-03 training 1.51e-02
   896 bs    128 per sample loss 5.56e-02 loading 5.56e-03 training 8.87e-03
  1792 bs    128 per sample loss 5.54e-02 loading 3.94e-03 training 5.89e-03
```

In [27]:

```
kubectll delete jobs --all || true
kubectll delete pods --all || true
```

```
job.batch "myjob" deleted
pod "myjob-fgflt" deleted
```

Kubernetes

- a way of running services and jobs on a cluster of machines
- configurations are given as JSON or YAML files (or via APIs)
- both CPUs and GPUs supported

In []: