

# Kubernetes Service Discovery

## Service Discovery

- for large scale deep learning we need multiple processes that talk to each other
- this requires
  - service discovery
  - networking
  - name resolution

## K8s Service Discovery

Simple:

- every pod gets assigned a hostname and domain
- you can simply connect directly to these well-known names

Requirements:

- create a "headless service" to start the name resolver
- add ports, host name, and subdomain to your pods

## Headless Service

The `clusterIP: None` makes it headless. (Other services are load balancing, which we don't want.)

In [28]:

```
kubectl apply -f - <<'EOF'
apiVersion: v1
kind: Service
metadata:
  name: bigdata19
spec:
  clusterIP: None
  ports:
    - port: 7880
      targetPort: 7880
  selector:
    app: bigdata19
EOF
```

service/bigdata19 unchanged

## A Visible Pod

This pod will be assigned the DNS name `shards.bigdata19`.

In [30]:

```
# nodes get assigned DNS names if they have a port and the app label matches the headless service

kubectl delete pod/shards || true
kubectl apply -f - <<'EOF'
apiVersion: v1
kind: Pod
metadata:
  name: shards
  labels:
    app: bigdata19
spec:
  containers:
  - name: shards
    image: gcr.io/research-191823/bigdata19
    command: ["serve-imagenet-shards", "-b", "96", "zpub://0.0.0.0:7880"]
    ports:
      - containerPort: 7880
  restartPolicy: Never
  hostname: shards
  subdomain: bigdata19
EOF
```

Error from server (NotFound): pods "shards" not found  
pod/shards created

In [ ]:

```
sleep 15
```

## DNS Debugging

In [32]:

```
# make sure resolution is working
kubectl exec -ti shards -- nslookup shards.bigdata19

# check resolv.conf file
kubectl exec -ti shards -- cat /etc/resolv.conf
```

Server: 10.64.0.10  
Address: 10.64.0.10#53

Name: shards.bigdata19.default.svc.cluster.local  
Address: 10.0.1.149

nameserver 10.64.0.10  
search default.svc.cluster.local svc.cluster.local cluster.local c.research-191823.internal google.i  
nternal  
options ndots:5

In [33]:

```
# check service running
kubectl get svc --namespace=kube-system

# check endpoints
kubectl get ep kube-dns --namespace=kube-system
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
default-http-backend	NodePort	10.64.0.106	<none>	80:31530/TCP	11h
heapster	ClusterIP	10.64.5.246	<none>	80/TCP	11h
kube-dns	ClusterIP	10.64.0.10	<none>	53/UDP,53/TCP	11h
metrics-server	ClusterIP	10.64.15.252	<none>	443/TCP	11h
NAME	ENDPOINTS			AGE	
kube-dns	10.0.0.130:53,10.0.2.130:53,10.0.0.130:53 + 1 more...			11h	

In [34]:

```
# when desperate, you can look through the kube-dns logs
kubectl get pods --namespace=kube-system -l k8s-app=kube-dns
kubectl get pods --namespace=kube-system -l k8s-app=kube-dns -o name | sed 1q |
while read pod; do
  kubectl logs --tail=3 --namespace=kube-system $pod -c kubedns
  kubectl logs --tail=3 --namespace=kube-system $pod -c dnsmasq
  kubectl logs --tail=3 --namespace=kube-system $pod -c sidecar
  kubectl logs --tail=3 --namespace=kube-system $pod -c prometheus-to-sd
done
```

NAME	READY	STATUS	RESTARTS	AGE
kube-dns-79868f54c5-h9m5c	4/4	Running	0	11h
kube-dns-79868f54c5-kh556	4/4	Running	0	11h
E1209 05:11:04.345848 1 reflector.go:201] k8s.io/dns/pkg/dns/dns.go:192: Failed to list *v1.Service: Get https://10.64.0.1:443/api/v1/services?resourceVersion=0: net/http: TLS handshake timeout				
E1209 05:11:04.346614 1 reflector.go:201] k8s.io/dns/pkg/dns/dns.go:189: Failed to list *v1.Endpoints: Get https://10.64.0.1:443/api/v1/endpoints?resourceVersion=0: net/http: TLS handshake timeout				
I1209 07:01:00.234994 1 dns.go:601] Could not find endpoints for service "bigdata19" in namespace "default". DNS records will be created once endpoints show up.				
I1209 05:04:01.718338 1 nanny.go:146] dnsmasq[23]: using nameserver 127.0.0.1#10053 for domain cluster.local				
I1209 05:04:01.718345 1 nanny.go:146] dnsmasq[23]: using nameserver 169.254.169.254#53				
I1209 05:04:01.718354 1 nanny.go:146] dnsmasq[23]: read /etc/hosts - 7 addresses				
I1209 05:04:02.521706 1 server.go:45] Starting server (options {DnsMasqPort:53 DnsMasqAddr:127.0.0.1 DnsMasqPollIntervalMs:5000 Probes:[{Label:kubedns Server:127.0.0.1:10053 Name:kubernetes.default.svc.cluster.local. Interval:5s Type:33} {Label:dnsmasq Server:127.0.0.1:53 Name:kubernetes.default.svc.cluster.local. Interval:5s Type:33}] PrometheusAddr:0.0.0.0 PrometheusPort:10054 PrometheusPath:/metrics PrometheusNamespace:kubedns})				
I1209 05:04:02.521760 1 dnsprobe.go:75] Starting dnsProbe {Label:kubedns Server:127.0.0.1:10053 Name:kubernetes.default.svc.cluster.local. Interval:5s Type:33}				
I1209 05:04:02.521898 1 dnsprobe.go:75] Starting dnsProbe {Label:dnsmasq Server:127.0.0.1:53 Name:kubernetes.default.svc.cluster.local. Interval:5s Type:33}				
I1209 05:04:03.749028 1 main.go:124] Taking source configs from kubernetes api server				
I1209 05:04:03.749033 1 main.go:85] Built the following source configs: [{kubedns localhost 10054 /metrics [probe_kubedns_latency_ms probe_kubedns_errors dnsmasq_misses dnsmasq_hits] 0xc4203e40a0}]				
I1209 05:04:03.749061 1 main.go:133] Running prometheus-to-sd, monitored target is kubedns localhost:10054				

In [35]:

```
sleep 15
```

## Logs of the Running Server

The server is chugging along nicely, sending out training batches to anybody who will listen.

In [36]:

```
kubectl logs shards | sed 10q
```

```
serving zpub://0.0.0.0:7880
0 rate 0.000000 msg/s throughput 0.00e+00 bytes/s
10 rate 5.538139 msg/s throughput 8.00e+07 bytes/s
20 rate 5.260502 msg/s throughput 7.60e+07 bytes/s
30 rate 5.167468 msg/s throughput 7.47e+07 bytes/s
40 rate 5.068958 msg/s throughput 7.33e+07 bytes/s
50 rate 5.030015 msg/s throughput 7.27e+07 bytes/s
60 rate 4.998792 msg/s throughput 7.22e+07 bytes/s
70 rate 5.002644 msg/s throughput 7.23e+07 bytes/s
80 rate 4.998627 msg/s throughput 7.22e+07 bytes/s
```

In [37]:

```
kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
shards	1/1	Running	0	4m32s

## Starting a Client

Here is a small network client that listens to training data and outputs statistics.

In [38]:

```
kubectl delete pod/client || true
kubectl apply -f - <<'EOF'
apiVersion: v1
kind: Pod
metadata:
  name: client
  labels:
    app: bigdata19
spec:
  containers:
  - name: client
    image: gcr.io/research-191823/bigdata19
    command: ["tensormon", "zsub://shards.bigdata19:7880"]
    stdin: true
    tty: true
  restartPolicy: Never
  hostname: client
  subdomain: bigdata19
EOF
```

Error from server (NotFound): pods "client" not found  
pod/client created

In [39]:

```
sleep 15
```

## Client Output

In [40]:

```
kubectl logs client | sed 10q
```

```
input: ['zsub://shards.bigdata19:7880']
zsub://shards.bigdata19:7880
connected
```

10	1.818 batches/s	174.525 samples/s	(batchsize: 96)
20	5.192 batches/s	498.429 samples/s	(batchsize: 96)
30	5.199 batches/s	499.126 samples/s	(batchsize: 96)
40	4.700 batches/s	451.197 samples/s	(batchsize: 96)
50	5.196 batches/s	498.807 samples/s	(batchsize: 96)

## Starting a DL Client on a GPU Node

In [41]:

```
kubectl delete job/myjob || true
kubectl apply -f - <<'EOF'
apiVersion: batch/v1
kind: Job
metadata:
  name: myjob
  labels:
    app: bigdata19
spec:
  backoffLimit: 0
  template:
    spec:
      containers:
        - name: myjob
          image: gcr.io/research-191823/bigdata19
          command:
            - "/bin/bash"
            - "-c"
            - |
              cp /files/*.py .
              python3 training.py --tensorcom zsub://shards.bigdata19:7880
          stdin: true
          tty: true
          resources:
            limits:
              nvidia.com/gpu: "1"
          volumeMounts:
            - mountPath: /files
              name: files
      nodeSelector:
        cloud.google.com/gke-accelerator: nvidia-tesla-t4
      restartPolicy: Never
      volumes:
        - configMap:
            name: files
          name: files
EOF
```

Error from server (NotFound): jobs.batch "myjob" not found  
job.batch/myjob created

In [42]:

```
sleep 10
```

In [43]:

```
kubectl logs job/myjob
```

```
/opt/conda/lib/python3.6/site-packages/torchvision/io/_video_opt.py:17: UserWarning: video reader based on ffmpeg c++ ops not available
  warnings.warn("video reader based on ffmpeg c++ ops not available")
Mon Dec  9 16:40:23 UTC 2019; myjob-q2d62; root; /workspace; GPU 0: Tesla T4 (UUID: GPU-fd29201b-d663-6697-b413-a761dceb23c8);
creating resnet50
 0 bs    96 per sample loss 7.45e-02 loading 6.03e-04 training 1.85e-02
```

In [44]:

```
sleep 60
```

## Training

- Note that with distributed preprocessing, loading is very fast.
- We will talk about the Tensorcom package later.

In [45]:

```
kubectll logs job/myjob
```

```
/opt/conda/lib/python3.6/site-packages/torchvision/io/_video_opt.py:17: UserWarning: video reader based on ffmpeg c++ ops not available
warnings.warn("video reader based on ffmpeg c++ ops not available")
Mon Dec 9 16:40:23 UTC 2019; myjob-q2d62; root; /workspace; GPU 0: Tesla T4 (UUID: GPU-fd29201b-d663-6697-b413-a761dceb23c8);
creating resnet50
    0 bs      96 per sample loss 7.45e-02 loading 6.03e-04 training 1.85e-02
   1152 bs    96 per sample loss 7.38e-02 loading 5.74e-04 training 7.39e-03
   2304 bs    96 per sample loss 7.34e-02 loading 5.62e-04 training 4.28e-03
   3456 bs    96 per sample loss 7.31e-02 loading 5.57e-04 training 3.42e-03
   4608 bs    96 per sample loss 7.29e-02 loading 5.56e-04 training 3.18e-03
   5760 bs    96 per sample loss 7.27e-02 loading 5.58e-04 training 3.13e-03
   6816 bs    96 per sample loss 7.26e-02 loading 5.56e-04 training 3.12e-03
```

In [46]:

```
kubectll get jobs
```

NAME	COMPLETIONS	DURATION	AGE
myjob	0/1	72s	72s

In [47]:

```
kubectll delete jobs --all
kubectll delete pods --all
```

```
job.batch "myjob" deleted
pod "client" deleted
pod "myjob-q2d62" deleted
pod "shards" deleted
```

## Kubernetes Service Discovery

- it's like creating new server out of thin air
- you can define your distributed application as a collection of pods
- K8s also provides load balancing and more complex name spaces